

# Retrieval Augmented Generation (RAG)

## Bridging Frozen Weights with Fresh Data

Iverina Ivanova & Gemini 3 Pro

Goethe Universität Frankfurt

January 12th, 2026

# Recap: The Frozen Encyclopedia

## Recall from Karpathy (Weeks 7-9):

- Training is like printing a massive encyclopedia.
- It is expensive and takes months.
- **Result:** The model's knowledge is **frozen** at the training date.

### 1. The Cutoff Issue

Model trained in 2024 doesn't know the 2025 election results.

### 2. The Specialization Issue

Model read the internet, but not *your* seminar paper.

# The Problem: The Confident Hallucination

We discussed **hallucinations** and the different strategies of mitigating them.

When an LLM doesn't know the answer, for example, because of its cutoff date, it relies on probability to guess the next word.

# The Problem: The Confident Hallucination

**Scenario:** You ask a model trained on data up to Dec. 2024:

*"Who is the current Chancellor of Germany?"*

## The Issue

The model does not know.

## The Reaction

It relies on probability.

- It might guess "**Angela Merkel**" (most frequent in training data).
- It might guess "**Olaf Scholz**" (last known state).
- **It cannot know the new reality without access to fresh data.**

# The Problem: The Confident Hallucination

**How do we fix this without retraining the whole model?**


# Strategies for mitigating hallucinations

- Tool calling (e.g., web search)
- Code execution

# Another strategy: RAG (Retrieval Augmented Generation)

**Concept:** Don't force the model to rely on its long-term memory. Give it access to external sources to consult before it answers the question.

## Standard LLM



*Closed-Book Exam*

Relies only on memorized facts.

## RAG



*Open-Book Exam*

Allowed to browse the textbook (your PDF) before answering.

# Why are we learning this now?

RAG brings the core topics we've covered so far together:

- ① **Vector Semantics:** How does the LLM find the right piece of information in a document to answer our question? It uses *embeddings* and *cosine similarity* to retrieve the relevant data.
- ② **LLM Generation:** Once the data is retrieved from the external sources, the generative LLM considers it before formulating its answer.
- ③ **Ethics & Reliability:** RAG helps ground AI in facts, reducing hallucinations and making sources traceable (citations).

# Under the Hood: Two Types of Memory

RAG combines two types of knowledge:

## 1. Parametric Knowledge (Internal Memory)

- What the model learned during pre-training.
- *Examples:* grammar, vocabulary, general facts (e.g., "Paris is in France").
- **Hard & expensive to update.**

## 2. Non-Parametric Knowledge (External Source)


- Your external data (PDFs, databases).
- The model looks this up in real-time.
- **Easy to update** (just add a new PDF).

# The Standard RAG Pipeline: 6 Steps

How do we build this system? It is a process of translation and search.  
Simple RAG Demo

- ① **Load:** Import your documents (PDFs, txt).
- ② **Split (Chunk):** Break text into small pieces.
- ③ **Embed:** Turn text into vectors.
- ④ **Store:** Save vectors in a database.
- ⑤ **Retrieve:** Find the most similar chunks to the user's question.
- ⑥ **Generate:** The LLM writes an answer using the retrieved chunks.

# How does RAG work?



Source: Langchain. URL:  
<https://python.langchain.com/docs/tutorials/rag/>

## Steps 1 & 2: Load and Split

- 1. Loading (Real-world data is messy) Challenge:** Computers see a stream of characters. Tables, footnotes, and multi-column layouts often break the reading process. If you feed garbage formatting into the RAG, you get garbage answers.
- 2. Splitting (The Chunking Dilemma) Challenge:** How big should a piece be?

### The Trade-Off: Context vs. Precision

- **Too Small (e.g., 1 sentence):** *Risk: Loss of Context.* (Example: A chunk says *He disagreed*. → The LLM doesn't know who *He* is or what he disagreed with.)
- **Too Large (e.g., 5 pages):** *Risk: Noise.* The retrieval might grab the right page, but the LLM gets distracted by 4 pages of irrelevant info surrounding the answer.

# Steps 1 & 2: Load and Split

**Goal:** Find the *goldilocks* zone: Enough text to make sense, small enough to be precise.

## Best practices

- fixed-size chunking (e.g., 512 tokens) with a chunk overlap (e.g., 50 tokens) that ensures context continuity.
- document-based chunking (e.g., using specialized libraries that segment the document intelligently based on its layout and data structures)

## Steps 3 & 4: Embed and Store

**Recall Session 5 (Vector Semantics):** Computers cannot *read* text; they understand numbers.

- We use an **embedding model** to turn our text chunks into lists of numbers (embeddings/vectors).
- We create indexes on the embeddings and store both the chunks and their corresponding embeddings in a **vector database**.

*"Friedrich Merz was elected chancellor on May 6th, 2025."* →  
[0.1, -0.5, 0.8, ...]


This allows us to search by **meaning**, not just by keyword matching.

## Step 5: The Mechanism (Cosine Similarity)

User Question: "Who is the current Chancellor of Germany?"

- ① **Vectorize:** Convert the question to a vector.  
[0.1, -0.9, 0.4...]
- ② **Compare:** Embeddings are represented as points in a vector space. Points which are clustered together are more similar in meaning.  
Small angle = High similarity.
- ③ **Retrieve Top-K:** Grab the top closest chunk to the question.  
*Match found: "Election results of 2025..."*

# Similarity search



# Step 6: Context Injection

We don't just ask the question. We build a **Composite Prompt**.

## The Hidden Prompt Structure

### 1. System Instruction:

*You are a helpful assistant. Answer using ONLY the provided context.*

### 2. Injected Context (The Retrieval):

*"...on January 15th, Friedrich Merz was elected..."*

### 3. User Question:

*"Who is the current Chancellor?"*

**Result:** The LLM generates a factual answer based on the injected context.

# Summary & Next Steps

## Why RAG matters:

- It allows us to chat with our own datasets without retraining the model.
- It provides citations (we know *where* the answer came from).
- It connects the frozen model to fresh data.

## Coming up:

- We'll discuss some RAG-related challenges.
- We'll explore RAG implementations:
  - **NotebookLM** (Google's RAG tool)
  - Code implementations with state-of-the-art Python libraries

# References

- Introduction to RAG in AI development. URL:  
<https://learn.microsoft.com/en-us/azure/databricks/generative-ai/tutorials/introduction-to-rag>
- Chunking Strategies for LLM Applications. URL:  
<https://www.pinecone.io/learn/chunking-strategies/>
- ChunkVisualizer. URL: <https://chunkviz.up.railway.app/explanation>
- Sentence Similarity. URL:  
<https://huggingface.co/tasks/sentence-similarity>

# Meta-Analysis

This slide deck was developed in a multi-step process involving **Gemini 3 Pro**.

- ① Providing the model with the necessary background of the course purpose and the target audience.
- ② Asking Gemini to explain RAG in simple terms and to relate its relevance to the previously discussed topics.
- ③ **Human Verification:** The specific definitions and workflows were cross-referenced with the technical documentations from the source pages to ensure factual accuracy.
- ④ Refining certain explanations by making them less wordy or by using visual representations.